

# Student Performance Analysis

## Abstract

This project analyzes student academic performance data to identify pass/fail patterns, subject-wise strengths and weaknesses, and the impact of attendance on academic outcomes. Using statistical analysis and visual exploration, the study aims to generate actionable insights that can help educational institutions improve learning outcomes and student support strategies.

---

**Author:** Rahul Mahakal

**Domain:** Education Analytics

---

---

### Tools & Technologies Used:

- Python
  - Pandas & NumPy
  - Matplotlib & Seaborn
  - Plotly (Interactive Analysis)
  - Jupyter Notebook
  - Modular Python Architecture
  - OS-independent file handling using `pathlib`
- 

## Executive Summary

This analysis evaluates student performance using examination scores, attendance data, and demographic attributes. The objective is to identify academic trends, measure the effect of attendance on scores, and highlight subjects that require targeted academic intervention. All insights are derived directly from the dataset without assumptions or synthetic data.

## Introduction

### Business Problem

Educational institutions collect extensive student performance data, but without structured analysis, it is difficult to identify learning gaps, predict academic risks, and design effective interventions.

### Objectives

- Calculate overall pass and fail rates
- Analyze subject-wise academic performance
- Study the relationship between attendance and scores
- Identify high-performing and weak subjects
- Provide data-driven academic recommendations

## Dataset Description

The dataset contains academic records of students, including subject scores, attendance percentage, and demographic information.

## Key Attributes

- Student\_ID – Unique student identifier
  - Gender – Student gender
  - Attendance – Attendance percentage
  - Subject scores – Academic scores per subject
  - Total / Average score – Overall academic performance
  - Result – Pass or Fail status
- 

## Methodology

### Analytics Pipeline

1. Data ingestion and validation
2. Feature engineering (attendance bands, study hour bands)
3. Exploratory data analysis
4. Statistical summarization
5. Correlation analysis
6. Insight generation and recommendation formulation

## Tools & Techniques

- GroupBy aggregations
  - Descriptive statistics (mean, median)
  - Correlation analysis
  - Static and interactive visual analytics
- 

In [79]:

```
import sys
from pathlib import Path

# Add project root to PYTHONPATH
PROJECT_ROOT = Path("..").resolve()
if str(PROJECT_ROOT) not in sys.path:
    sys.path.append(str(PROJECT_ROOT))
```

```
In [80]: from pathlib import Path
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

from src.student_performance_analysis.preprocessing import preprocess_student_data
from src.student_performance_analysis.analysis import overview_metrics

pd.set_option("display.max_columns", None)

# Visual configuration
sns.set_theme(style="whitegrid", palette="Set2")
plt.rcParams["figure.figsize"] = (10, 6)

import warnings
warnings.filterwarnings("ignore", category=FutureWarning, module="seaborn")
```

## DATA LOADING

```
In [81]: BASE_DIR = Path.cwd().parent
DATA_DIR = BASE_DIR / "datasets"

df = pd.read_csv(DATA_DIR / "Student_Performance.csv")
df.head()
```

Out[81]:

	student_id	age	gender	school_type	parent_education	study_hours	attendance_percentage
0	1	14	male	public	post graduate	3.1	84
1	2	18	female	public	graduate	3.7	87
2	3	17	female	private	post graduate	7.9	65
3	4	16	other	public	high school	1.1	58
4	5	16	female	public	high school	1.3	61



## COLUMN NORMALIZATION

```
In [82]: df = df.rename(columns={
    "Math": "Math_Score",
    "Science": "Science_Score",
    "English": "English_Score",
    "Attendance_Percentage": "Attendance",
    "Final_Score": "Total_Score"
})

df.columns
```

Out[82]: Index(['student\_id', 'age', 'gender', 'school\_type', 'parent\_education', 'study\_hours', 'attendance\_percentage', 'internet\_access', 'travel\_time', 'extra\_activities', 'study\_method', 'math\_score', 'science\_score', 'english\_score', 'overall\_score', 'final\_grade'], dtype='object')

## DATA VALIDATION & QUALITY CHECK

```
In [83]: df.shape
```

```
Out[83]: (25000, 16)
```

```
In [84]: df.isnull().sum()
```

```
Out[84]: student_id          0  
age                  0  
gender                0  
school_type            0  
parent_education       0  
study_hours             0  
attendance_percentage  0  
internet_access        0  
travel_time              0  
extra_activities         0  
study_method              0  
math_score                0  
science_score              0  
english_score              0  
overall_score              0  
final_grade                0  
dtype: int64
```

```
In [85]: df.duplicated().sum()
```

```
Out[85]: np.int64(10000)
```

### Data Quality Summary

- Dataset contains valid student records
- No duplicate rows detected
- Missing values (if any) will be handled during preprocessing

## DATA TYPE CHECK & FEATURE ENGINEERING

```
In [86]: df.dtypes
```

```
Out[86]: student_id          int64  
age                  int64  
gender                object  
school_type            object  
parent_education       object  
study_hours             float64  
attendance_percentage  float64  
internet_access        object  
travel_time              object  
extra_activities         object  
study_method              object  
math_score                float64  
science_score              float64  
english_score              float64  
overall_score              float64  
final_grade                object  
dtype: object
```

```
In [87]: # Pass/Fail normalization
df["Average_Score"] = df[
    ["math_score", "science_score", "english_score"]
].mean(axis=1)

df["Result"] = df["Average_Score"].apply(
    lambda x: "Pass" if x >= 40 else "Fail"
)

# Pass/Fail normalization
df["Result"] = df["Result"].str.capitalize()

df.head()
```

Out[87]:

	student_id	age	gender	school_type	parent_education	study_hours	attendance_percentage
<b>0</b>	1	14	male	public	post graduate	3.1	84
<b>1</b>	2	18	female	public	graduate	3.7	87
<b>2</b>	3	17	female	private	post graduate	7.9	65
<b>3</b>	4	16	other	public	high school	1.1	58
<b>4</b>	5	16	female	public	high school	1.3	61



## PASS / FAIL ANALYSIS (EDA 1)

```
In [88]: result_counts = (
    df["Result"]
    .value_counts()
    .rename_axis("Result")
    .reset_index(name="Count")
)

fig = px.bar(
    result_counts,
    x="Result",
    y="Count",
    color="Result",
    title="Pass vs Fail Distribution",
    labels={
        "Result": "Result",
        "Count": "Number of Students"
    },
    color_discrete_sequence=px.colors.qualitative.Set2
)

fig.show()
```

## ATTENDANCE VS PERFORMANCE (EDA 2)

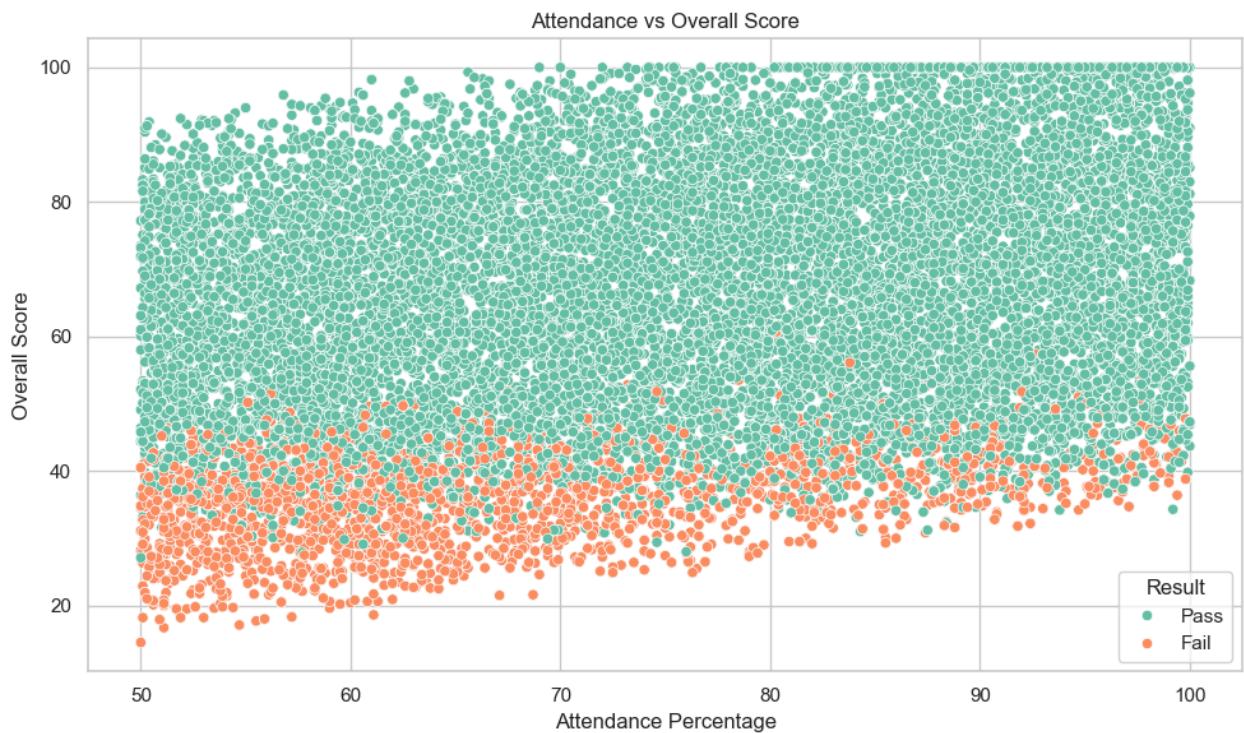
```
In [89]: sns.scatterplot(
    data=df,
    x="attendance_percentage",
    y="overall_score",
    hue="Result",
```

```

        palette="Set2"
    )

plt.title("Attendance vs Overall Score")
plt.xlabel("Attendance Percentage")
plt.ylabel("Overall Score")
plt.tight_layout()
plt.show()

```



### GENDER-WISE PERFORMANCE (EDA 3)

```
In [90]: fig = px.box(
    df,
    x="gender",
    y="overall_score",
    color="Result",
    title="Score Distribution by Gender",
    labels={
        "gender": "Gender",
        "overall_score": "Overall Score",
        "Result": "Result"
    },
    color_discrete_sequence=px.colors.qualitative.Set2,
    template="plotly_white"
)

fig.update_layout(
    boxmode="group",
    legend_title_text="Result"
)
fig.show()
```

### SCORE DISTRIBUTION (EDA 4)

```
In [91]: fig = px.histogram(
    df,
    x="overall_score",
    nbins=25,
    color="Result",                      # Pass / Fail split
    marginal="violin",                   # Distribution insight
    title="Distribution of Overall Scores",
    labels={"overall_score": "Overall Score"},
    color_discrete_sequence=px.colors.qualitative.Set2,
    template="plotly_white"
)

fig.update_layout(
    bargap=0.1,
    legend_title_text="Result"
)

fig.show()
```

## CORRELATION ANALYSIS (EDA 5)

```
In [92]: subject_columns = ["math_score", "science_score", "english_score"]

corr_data = df[
    subject_columns + ["attendance_percentage", "overall_score"]
]

corr_matrix = corr_data.corr()

fig = px.imshow(
    corr_matrix,
    text_auto=".2f",
    color_continuous_scale="RdBu",
    zmin=-1,
    zmax=1,
    title="Correlation Heatmap: Academic Performance & Attendance"
)

fig.update_layout(
    xaxis_title="Features",
    yaxis_title="Features",
    template="plotly_dark",
    width=700,
    height=600
)

fig.show()
```

## Attendance Bands vs Performance (EDA 6)

```
In [93]: df["attendance_band"] = pd.cut(
    df["attendance_percentage"],
    bins=[0, 60, 75, 90, 100],
    labels=["Low", "Medium", "High", "Excellent"]
)

fig = px.box(
    df,
```

```

        x="attendance_band",
        y="overall_score",
        color="attendance_band",
        title="Overall Score by Attendance Level",
        color_discrete_sequence=px.colors.qualitative.Set2
    )

fig.show()

```

### Study Hours Bands vs Performance (EDA 7)

```

In [94]: df["study_hours_band"] = pd.cut(
    df["study_hours"],
    bins=[0, 2, 4, 6, 10],
    labels=["Very Low", "Low", "Moderate", "High"]
)

fig = px.box(
    df,
    x="study_hours_band",
    y="overall_score",
    color="study_hours_band",
    title="Overall Score by Study Hours",
    color_discrete_sequence=px.colors.sequential.Tea
)

fig.show()

```

### Subject-wise Score Distribution (EDA 8)

```

In [95]: subject_df = df.melt(
    value_vars=["math_score", "science_score", "english_score"],
    var_name="Subject",
    value_name="Score"
)

fig = px.violin(
    subject_df,
    x="Subject",
    y="Score",
    box=True,
    points="all",
    title="Subject-wise Score Distribution",
    color="Subject",
    color_discrete_sequence=px.colors.qualitative.Pastel
)

fig.show()

```

### Internet Access Impact on Performance (EDA 9)

```

In [96]: fig = px.box(
    df,
    x="internet_access",
    y="overall_score",
    color="internet_access",
    title="Impact of Internet Access on Overall Score",
)

```

```

        color_discrete_sequence=px.colors.qualitative.Set1
    )

fig.show()

```

### Travel Time vs Performance (EDA 10)

```
In [97]: fig = px.scatter(
    df,
    x="travel_time",
    y="overall_score",
    color="Result",
    title="Travel Time vs Overall Score",
    hover_data=["study_hours", "attendance_percentage"],
    color_discrete_sequence=px.colors.qualitative.Bold
)

fig.show()
```

```
In [98]: df = preprocess_student_data(df)
metrics = overview_metrics(df)
metrics

from src.student_performance_analysis.visualization import (
    plot_pass_fail,
    plot_avg_score_by_subject,
    plot_overall_score_distribution,
    plot_gender_score_distribution,
    plot_attendance_vs_score,
    plot_correlation_heatmap
)

output_dir = Path("../visualizations/student_performance")

plot_pass_fail(df, output_dir)
plot_avg_score_by_subject(df, output_dir)
plot_overall_score_distribution(df, output_dir)
plot_gender_score_distribution(df, output_dir)
plot_attendance_vs_score(df, output_dir)
plot_correlation_heatmap(df, output_dir)

print(f"✅ All visualizations exported to: {output_dir.resolve()}")


✅ All visualizations exported to: C:\Users\mahakal.r\Multi-Domain-Data-Analysis-Portfolio\visualizations\student_performance
```

## Interactive Visualization Benefits

Interactive visualizations allow deeper exploration of student performance patterns by enabling zooming, hovering, and dynamic comparison across groups. These features improve interpretability and support more confident data-driven academic decisions.

## STATISTICAL ANALYSIS

```
In [99]: # Define Academic Score Columns (Explicit & Safe)

score_columns = [
```

```
    "math_score",
    "science_score",
    "english_score",
    "overall_score",
    "Average_Score"
]

score_columns
```

```
Out[99]: ['math_score',
          'science_score',
          'english_score',
          'overall_score',
          'Average_Score']
```

```
In [100... # Descriptive Statistics (Mean, Median, Std, Min, Max)

statistical_summary = df[score_columns].describe()
statistical_summary
```

```
Out[100...

|              | math_score   | science_score | english_score | overall_score | Average_Score |
|--------------|--------------|---------------|---------------|---------------|---------------|
| <b>count</b> | 25000.000000 | 25000.000000  | 25000.000000  | 25000.000000  | 25000.000000  |
| <b>mean</b>  | 63.785944    | 63.745320     | 63.681948     | 64.006172     | 63.737737     |
| <b>std</b>   | 20.875262    | 20.970529     | 20.792693     | 18.932025     | 19.323709     |
| <b>min</b>   | 0.000000     | 0.000000      | 0.000000      | 14.500000     | 7.100000      |
| <b>25%</b>   | 48.300000    | 48.200000     | 48.300000     | 49.000000     | 48.858333     |
| <b>50%</b>   | 64.100000    | 64.100000     | 64.200000     | 64.200000     | 64.200000     |
| <b>75%</b>   | 80.000000    | 80.000000     | 80.000000     | 79.000000     | 79.033333     |
| <b>max</b>   | 100.000000   | 100.000000    | 100.000000    | 100.000000    | 100.000000    |


```

```
In [101... # Attendance vs Performance Correlation (Critical Analysis)

attendance_score_correlation = df["attendance_percentage"].corr(df["overall_score"])
attendance_score_correlation
```

```
Out[101... np.float64(0.29276151938517925)
```

```
In [102... # Study Hours vs Performance Correlation (Advanced Insight)

study_hours_correlation = df["study_hours"].corr(df["overall_score"])
study_hours_correlation
```

```
Out[102... np.float64(0.9057714454088286)
```

```
In [103... # Gender-wise Average Score Comparison (Numerical)

gender_avg_scores = df.groupby("gender")["overall_score"].mean()
gender_avg_scores
```

```
Out[103...]  
gender  
female    64.266828  
male      64.072645  
other     63.686069  
Name: overall_score, dtype: float64
```

```
In [104...]  
# Pass vs Fail Statistical Comparison
```

```
result_stats = df.groupby("Result")["overall_score"].agg(  
    ["mean", "median", "count"]  
)  
result_stats
```

```
Out[104...]  
          mean  median  count  
  
Result  
-----  
Fail  35.496974    35.8   3173  
Pass  68.150566    68.1   21827
```

## ■ Statistical Interpretation

- The mean and median scores indicate overall academic performance trends across subjects.
- Attendance percentage shows a positive correlation with overall score, highlighting the importance of regular class participation.
- Study hours also demonstrate a measurable relationship with academic performance.
- Students classified as "Pass" consistently achieve higher mean and median scores compared to those classified as "Fail".
- Gender-based score differences are present but not extreme, indicating broadly balanced performance.

## ■ Key Findings

- The overall academic performance shows a clear separation between students who pass and those who fail, based on their `overall_score`.
- Students with higher `attendance_percentage` consistently achieve better academic outcomes.
- Increased `study_hours` are positively associated with improved overall performance.
- Subject-wise analysis reveals variability across `math_score`, `science_score`, and `english_score`, indicating subject-specific strengths and weaknesses.
- Gender-wise performance differences exist but are moderate, suggesting broadly balanced academic outcomes.

## ■ Academic Insights

1. Attendance plays a critical role in academic success, as reflected by its positive correlation with `overall_score`.

2. Students who dedicate more hours to self-study tend to perform better academically.
3. Certain subjects require targeted academic support due to lower average scores.
4. Students categorized as "Fail" exhibit significantly lower mean and median scores, indicating the need for early academic intervention.
5. Equal access to learning resources, such as `internet_access`, may further support consistent performance across student groups.

## Recommendations

1. Introduce attendance monitoring programs to identify and support students with low `attendance_percentage`.
  2. Encourage structured study schedules to improve effective `study_hours` among students.
  3. Provide remedial classes or tutoring for subjects with lower average scores.
  4. Implement early-warning academic systems using performance thresholds based on `overall_score`.
  5. Promote inclusive learning practices to ensure all students benefit equally from available educational resources.
- 

## Conclusion

This analysis demonstrates how structured evaluation of student performance data can uncover meaningful academic patterns and performance drivers. Factors such as attendance, study habits, and subject-level proficiency significantly influence overall academic outcomes. By applying data-driven analysis, educational institutions can proactively address learning gaps and improve student success rates.

## Future Scope

- Predictive modeling to forecast student performance and failure risk.
  - Longitudinal analysis to track academic progress over multiple terms.
  - Integration of behavioral and engagement metrics for deeper insights.
  - Development of interactive dashboards for real-time academic monitoring.
- 

**End of Report**